

Range Minimum Queries

Feyzi Eren Gündoğdu

Karlsruhe Institute of Technology, Germany

1 Introduction

This report presents a comparison of different data structures and their performance for Range Minimum Queries (RMQ) in C++.

Six different data structures, including two naive and four advanced approaches, are discussed in this report. The naive algorithms, *OnTheFlyNaive* and *PrecomputedNaive*, are included to provide upper-bound performance comparisons for RMQ. The advanced data structures considered are *SparseTable*, *SegmentTree*, *BlockSparseTable*, and *CartesianTree*.

For some of the advanced data structures, different variants were evaluated and are represented by different suffixes (e.g., *BST16* denotes a Block Sparse Table with a block size of 16). n represents the input size, B represents block size.

TODO FILES

2 Methods

Algorithm	Space overhead (bits)	Query time
OnTheFlyNaive	$O(1)$	$O(n)$
PrecomputedNaive	$O(n^2 \log n)$	$O(1)$
Sparse Table	$O(n \log n)$	$O(1)$
Segment Tree	$O(n \log n)$	$O(\log n)$
Blocks	$O\left(\frac{n}{s} \log^2 n\right)$	$O(s)$
Cartesian Trees	$O(n \log n)$	$O(1)$

Table 1 Theoretical bounds.

2.1 OnTheFlyNaive

This algorithm stores only a pointer to the input data and answers each query by scanning the requested range on the fly. Since no preprocessing is performed and no additional RMQ data structure is stored, it requires only $O(1)$ extra space. In the worst case, a query may scan the whole array, so the query time is $O(n)$.

2.2 PrecomputedNaive

This algorithm stores the answers for all possible RMQ instances in a triangular table. For every left endpoint l , it stores the minimum index of each range $[l, r]$ with $l \leq r$. Since only ranges with $l \leq r$ are valid, it is sufficient to store the upper triangular part of the full $n \times n$ table. In the implementation, the entry for the range $[l, r]$ is stored at position `table[l][r - l]`. Equivalently, in a flat triangular layout, the range $[l, r]$ can be mapped to a single array position by

$$l \cdot n - \frac{l(l-1)}{2} + (r-l).$$

A query can then be answered by directly accessing the corresponding table entry. The algorithm theoretically requires $O(n^2 \log n)$ bits of space and query time is $O(1)$. For this

implementation, each entry is stored as a fixed-width `uint32_t` instead of using exactly $\lceil \log_2 n \rceil$ bits per entry. Therefore, the corresponding theoretical space complexity is

$$O\left(\frac{n(n+1)}{2} \log n\right) = O(n^2 \log n).$$

2.3 SparseTable

This algorithm stores a sparse table as a nested vector with $k + 1$ levels, where $k = \lfloor \log_2 n \rfloor$. At level ℓ , the table stores one entry for each valid interval of length 2^ℓ , resulting in $n + 1 - 2^\ell$ entries. Each entry stores the index of the minimum element in the corresponding interval. In the implementation, these indices are stored using `size_t`. To answer a query $[l, r]$, the algorithm chooses $\ell = \lfloor \log_2(r - l + 1) \rfloor$ and combines the two overlapping intervals of length 2^ℓ that cover the query range. This requires only two table accesses and one comparison, so the theoretical query time is $O(1)$. The exact number of stored entries is

$$\sum_{\ell=0}^{\lfloor \log_2 n \rfloor} (n + 1 - 2^\ell).$$

Thus, the sparse table stores $O(n \log n)$ entries. Since each entry is stored as a fixed-size `size_t` value in the implementation, the theoretical space usage is proportional to $O(n \log n)$.

3 Results

3.1 OnTheFlyNaive

3.1.1 Space

In the bits-per-element plot in Figure 1, the space usage appears to converge to 0. This is expected, because the method uses only constant additional space. In this implementation, the only stored data is a pointer to the input array, which takes approximately 8 bytes. Therefore, the plotted value is approximately $64/n$ bits per element, which tends to 0 as n grows. The absolute space overhead is shown in Figure 2. There, the space usage remains constant at approximately 64 bits, matching the theoretical $O(1)$ extra space bound.

3.1.2 Time

The query time in Figure 3 grows with n , which is consistent with the $O(n)$ worst-case query time. This is because each query is answered by scanning the queried range directly.

3.1.3 Tradeoff

3.2 PrecomputedNaive

3.2.1 Space

The results from `data.csv` indicate that the measured space usage matches the expected implementation-level bound. For example, for $n = 10000$, the table contains

$$\frac{10000 \cdot 10001}{2}$$

entries. Since each entry is stored as a `uint32_t`, each entry occupies 4 bytes. Therefore, the expected space usage is approximately

$$4 \cdot \frac{10000 \cdot 10001}{2} = 200020000$$

bytes, plus a small constant overhead. This matches the measured value 200020040 bytes. Although the curve appears almost linear in the log-log Figure 1, this does not indicate linear space growth in n . The measured space usage of `PrecomputedNaive` is proportional to

$$32 \cdot \frac{n(n+1)}{2}$$

bits, which is $\Theta(n^2)$. On a log-log plot, polynomial growth appears as a straight line; in this case, the slope corresponds to quadratic growth.

3.2.2 Time

For the query time, the theoretical $O(1)$ bound is only approximately visible in the measurements. Each query performs a constant number of table lookups, but the measured running time is affected by cache and memory hierarchy effects as the table grows. Nevertheless, Figure 3 shows that `PrecomputedNaive` behaves near-constantly. The query time of `PrecomputedNaive` increases sharply from $n = 1000$ to $n = 3000$, but grows more slowly afterwards. This can be explained by cache effects. For small inputs, a large part of the precomputed table can still be served from cache. Once the table exceeds the relevant cache levels, random table lookups become much more expensive. After this point, the queries are already mostly memory-latency bound, so increasing n further does not necessarily lead to a proportional increase in query time. Thus, the observed behavior is consistent with an $O(1)$ query algorithm affected by the memory hierarchy.

3.2.3 Tradeoff

3.3 SparseTable

3.3.1 Space

Numerically, the measurements from `data.csv` follow the theoretical bounds discussed above. In Figure 2, the absolute space overhead of the sparse table increases slightly faster than linearly. This is expected, since the sparse table stores entries for all power-of-two interval lengths, giving an $O(n \log n)$ space usage. The behavior in Figure 1 is also consistent with the theory. Since this plot shows the space usage normalized by the input size, the expected growth for the sparse table is

$$\frac{O(n \log n)}{n} = O(\log n).$$

Thus, the slow increase in bits per element matches the expected logarithmic growth.


3.3.2 Time

For the query time analysis, Figure 3 shows that the sparse table has almost constant query time for the smaller input sizes, which is consistent with the theoretical $O(1)$ query bound. For larger input sizes, however, the measured query time increases. This does not contradict the theoretical complexity. A sparse table query still performs only two table accesses and

4 Range Minimum Queries

one comparison. The increase in running time is caused by practical hardware effects. As the input size grows, the sparse table becomes too large to fit into the CPU caches. Consequently, the two table accesses are more likely to cause cache misses and main-memory accesses, which have higher latency. Thus, the measured running time increases due to memory hierarchy effects.

3.3.3 Tradeoff



■ **Figure 1** space usage




Figure 2 space overhead





Figure 3 query time

6 Range Minimum Queries



■ **Figure 4** space-time trade-off